

Homework 2: Decision Trees, Logistic Regression, and Support Vector Machines (100 points)

Prof. Amine Aboussalah

Due Date: Monday 3rd November 2025.

Submission will be done on Brightspace.

Datasets

- **Decision Trees:** [UCI Car Evaluation Dataset](#)
- **Logistic Regression:** [UCI Breast Cancer Wisconsin Dataset](#)
- **Support Vector Machines:** [UCI Breast Cancer Wisconsin Dataset](#)

Overview

In this homework, you'll explore support vector machines, decision trees, and logistic regression and apply them to real-world classification tasks.

You must **implement all algorithms from scratch** without using Scikit-learn or similar libraries. If your implementation runs slowly, you are permitted to reduce the number of datapoints or/and features used from the provided datasets to ensure reasonable execution time, **but you have to mention it.**

Task 1: Decision Tree Classifier (30 points)

Objective

Implement a decision tree classifier from scratch using information gain and Gini impurity criteria.

Instructions

Implementation (20 points)

- Implement the decision tree algorithm with support for:
 - Gini impurity and information gain criteria (5 points)
 - Handling both categorical and numerical features (5 points)
 - Recursive tree construction with depth and minimum split size constraints (5 points)
 - Prediction on test samples (5 points)

Application & Evaluation (10 points)

- Load and preprocess the **Car Evaluation dataset** (e.g., convert categorical features to numerical form) (2 points)
- Split the dataset into 80% training and 20% test (2 points)
- Train your decision tree with different criteria (gini vs information gain) (2 points)
- Evaluate performance using:
 - Accuracy, precision, recall, F1-score (1 point)
 - Confusion matrix (1 point)
- Compare your implementation to scikit-learn's built-in tree classifier. Analyze and interpret the obtained results. (2 points)

Task 2: Logistic Regression (30 points)

Objective

Implement logistic regression from scratch using gradient descent, and apply it to a binary classification problem.

Instructions

Implementation (20 points)

- Implement logistic regression for binary classification using sigmoid activation. (8 points)

- Add support for:
 - Mini-batch gradient descent (4points)
 - L2 regularization (4 points)
 - Early stopping based on validation loss (4 points)

Application & Evaluation (10 points)

- Load and preprocess the UCI Breast Cancer Wisconsin Dataset (e.g., handle missing values, normalize) (2 points)
- Use 70/15/15 train/validation/test split (2 points)
- Evaluate with accuracy, precision, recall, F1-score, and ROC-AUC (3 points)
- Compare your model to [scikit-learn's LogisticRegression](#) (3 points)

Task 3: Support Vector Machine Implementation (40 points)

Objective

Implement an SVM classifier from scratch and apply it to a binary classification problem.

Instructions

1. **Implementation** (20 points)
 - Code the SVM algorithm from scratch without using existing SVM libraries
 - Your implementation must include:
 - Core SVM algorithm design (8 points)
 - Linear kernel implementation (4 points)
 - Support for at least one non-linear kernel (RBF preferred) (4 points)
 - Optimization using Scipy library and prediction functionality (4 points)
2. **Application & Evaluation** (12 points)
 - Apply your SVM to the provided binary classification dataset (3 points)
 - Split the data into training (70%) and test (30%) sets (2 points)

- Train your SVM with different hyperparameters: kernel type (e.g. linear, polynomial, RBF ...) , kernel parameters (e.g. variance in the RBF kernel) (3 points)
 - Evaluate performance using: (4 points)
 - Accuracy, precision, recall, F1-score
 - ROC curve and AUC
 - Confusion matrix
3. **Analysis** (8 points)
- Analyze how different hyperparameters affect decision boundaries (3 points)
 - Visualize support vectors and decision boundaries for a 2D projection of your data (3 points)
 - Compare your implementation with scikit-learn's SVM (performance and efficiency), and previous logistic regression results (2 points)

General Requirements (up to -10 points penalty)

- Submit as a **Jupyter notebook**
- All code must be original
- Document each step and algorithm clearly
- Code must include comments and basic error handling
- Clearly report and explain all hyperparameter choices